

# WHEN DO CURRICULA WORK?

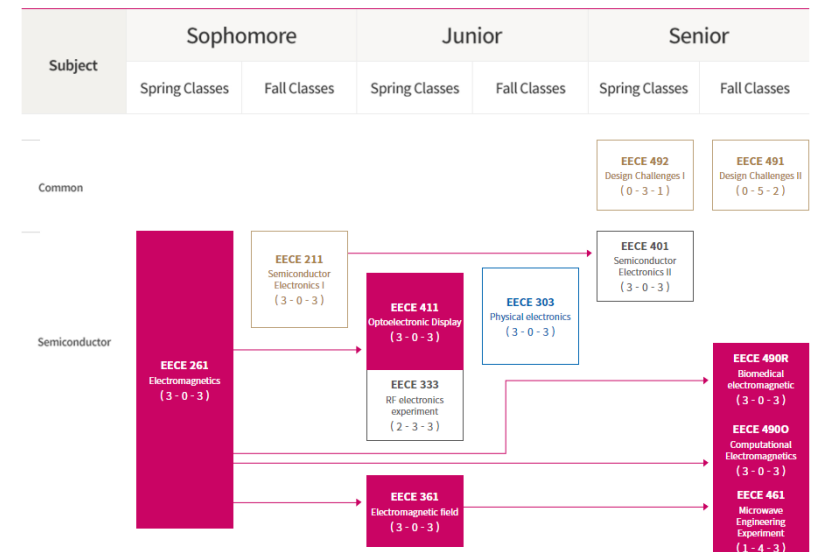
Xiaoxia Wu, Ethan Dyer, Behnam Neyshabur

Efficient Learning Lab@POSTECH

Junwon Seo

# Curriculum Learning

- Inspired by the importance of the **ordered** method when teaching humans
- Propose training models by presenting **easier** examples **earlier** during training



# Generality of Curriculum Learning

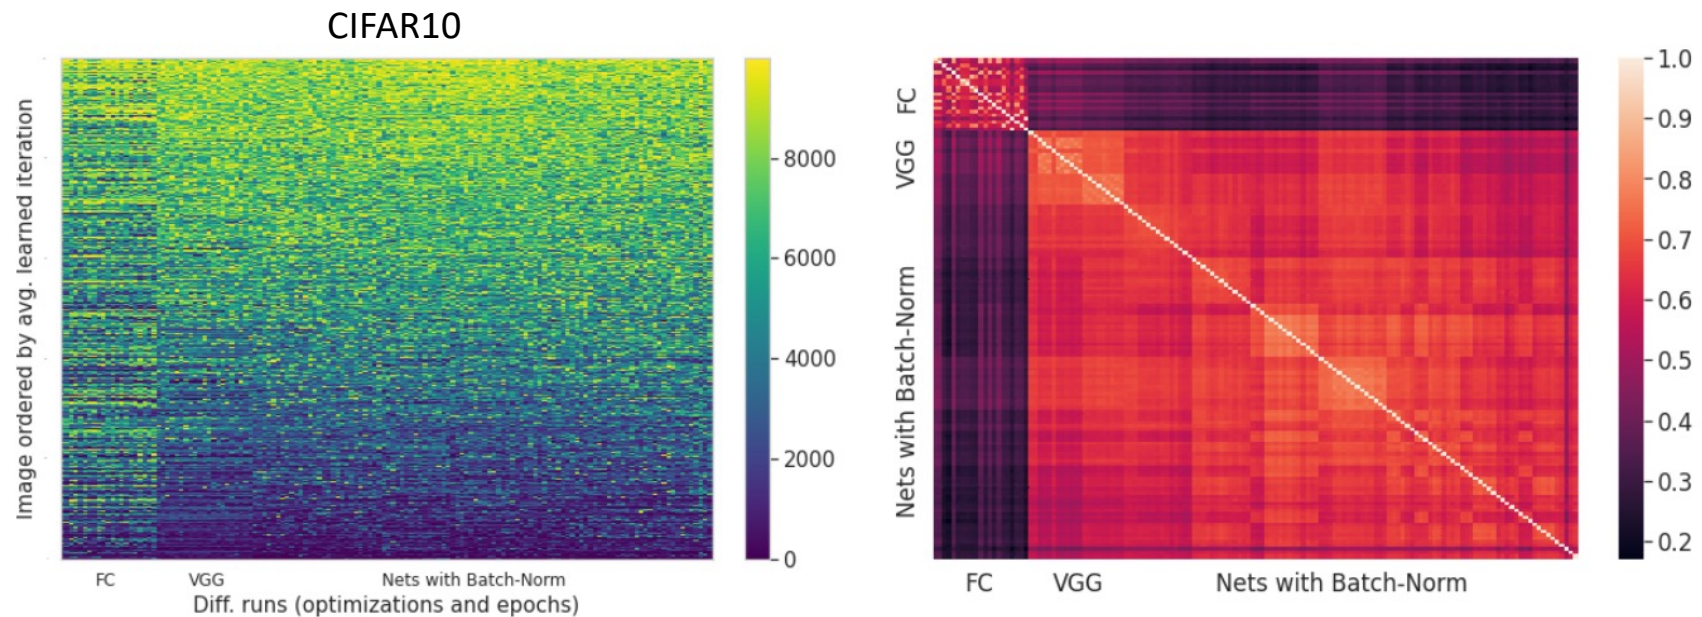
- Anti-curriculum learning
  - Can be as good as curriculum learning in certain scenarios.
  - Underperform standard or curriculum learning in other contexts.
- Empirical observations on curricula appear to be in conflict
- Despite a rich literature(about 50), do we use it?
  - No ordered learning method is known to improve consistently across contexts.

# *Implicit Curricula*

- The order in which a network learns examples under traditional stochastic gradient descent with i.i.d. data sampling
- How to quantify the order?
  - Define the *learned iteration* of a sample for a given model
  - The epoch for which the model correctly predicts the sample and all subsequent epochs
  - Explicitly,

$$\min_{t^*} \{t^* | \hat{y}_{\mathbf{w}}(t)_i = y_i, \forall t^* \leq t \leq T\}$$

# ***Implicit Curricula***



- FC/VGG/ResNet/WideResNet/DenseNet/EfficientNet B0/VGG-BN and Adam/SGD with momentum
- “At least within model types, **less ambiguity** about the difficulty of **a given image**”

# Three Ingredients in Curriculum learning

- The scoring function
  - A map from an input example,  $x$ , to a numerical score  $s(x) \in \mathbb{R}$
  - Higher score corresponds to a more difficult example
- The pacing function
  - The size of the training dataset used at each step  $t$ .
- The order
  - Curriculum: from **lowest** score to **highest** score
  - Anti-curriculum: from **highest** score to **lowest** score
  - Random

# Three Ingredients in Curriculum learning

---

**Algorithm 1** (Random-/Anti-) Curriculum learning with pacing and scoring functions


---

- 1: **Input:** Initial weights  $\mathbf{w}^0$ , training set  $\{\mathbf{x}_1, \dots, \mathbf{x}_N\}$ , pacing function  $g : [T] \rightarrow [N]$ , scoring function  $s : [N] \rightarrow \mathbb{R}$ , order  $o \in \{\text{“ascending”}, \text{“descending”}, \text{“random”}\}$ .
  - 2:  $(\mathbf{x}_1, \dots, \mathbf{x}_N) \leftarrow \text{sort}(\{\mathbf{x}_1, \dots, \mathbf{x}_N\}, s, o)$
  - 3: **for**  $t = 1, \dots, T$  **do**
  - 4:    $\mathbf{w}^{(t)} \leftarrow \text{train-one-epoch}(\mathbf{w}^{(t-1)}, \{\mathbf{x}_1, \dots, \mathbf{x}_{g(t)}\})$
  - 5: **end for**
- 

**Random ordering** corresponds to **i.i.d. training** on a training dataset with **dynamic size**

| Name       | Expression $g_{(a,b)}(t)$  |
|------------|--|
| log        | $Nb + N(1-b) \left(1 + .1 \log \left(\frac{t}{aT} + e^{-10}\right)\right)$       |
| exp        | $Nb + \frac{N(1-b)}{e^{10}-1} \left(\exp \left(\frac{10t}{aT}\right) - 1\right)$ |
| step       | $Nb + N \left\lceil \frac{x}{aT} \right\rceil$                                   |
| polynomial | $Nb + N \frac{(1-b)}{(aT)^p} t^p \quad - p = 1/2, 1, 2$                          |

Pacing Functions



a: the fraction of training needed to reach the size of the full training set

b: the fraction of the training set used at the start of training

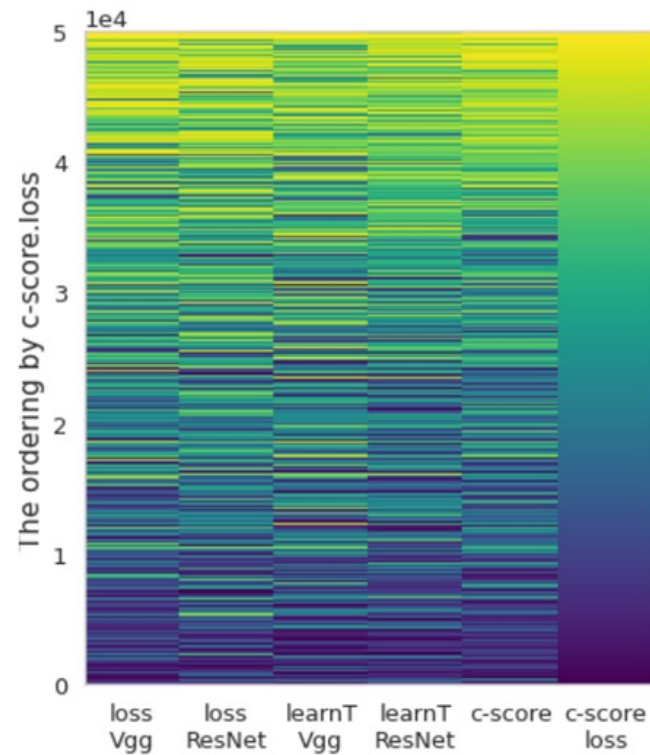
# Three Ingredients: Scoring Functions

- Scoring Function

- Loss function  $s(\mathbf{x}_i, y_i) = \ell(f_{\mathbf{w}}(\mathbf{x}_i), y_i)$ .
  - Samples are scored using the real-valued loss of a reference model
- Learned epoch/iteration  $s(\mathbf{x}_i, y_i) = \min_{t^*} \{t^* | \hat{y}_{\mathbf{w}}(t)_i = y_i, \forall t^* \leq t \leq T\}$ 
  - The epoch/iteration for which the model correctly predicts the sample and all subsequent epochs
- Estimated c-score  $s(\mathbf{x}_i, y_i) = \mathbb{E}_{D \sim \hat{\mathcal{D}} \setminus \{(\mathbf{x}_i, y_i)\}} [\mathbb{P}(\hat{y}_{\mathbf{w}, i} = y_i | D)]$  where  $D$ , with  $|D| = n$ 
  - consistency of a reference model correctly predicting a particular example's label when trained on independent i.i.d. draws of a fixed-sized dataset not containing that example.



# Three Ingredients: Scoring Functions



CIFAR10 (epoch 10)

Correlation Matrix

|               |      |      |      |      |      |      |
|---------------|------|------|------|------|------|------|
| loss Vgg      | 1.00 | 0.63 | 0.66 | 0.62 | 0.76 | 0.73 |
| loss ResNet   | 0.63 | 1.00 | 0.56 | 0.59 | 0.70 | 0.69 |
| learnT Vgg    | 0.66 | 0.56 | 1.00 | 0.79 | 0.69 | 0.71 |
| learnT ResNet | 0.62 | 0.59 | 0.79 | 1.00 | 0.72 | 0.80 |
| c-score       | 0.76 | 0.70 | 0.69 | 0.72 | 1.00 | 0.82 |
| c-score loss  | 0.73 | 0.69 | 0.71 | 0.80 | 0.82 | 1.00 |

loss Vgg loss ResNet learnT Vgg learnT ResNet c-score c-score loss


Epoch: {2, 10, 30, 60, 90, 120, 150, 180, 200}

ResNet-18(200 epoch) is only outlier


**DIFFICULTY SCORES ARE BROADLY CONSISTENT**

# Three ingredients: Pacing Functions

“Does curriculum learning learn the dataset in the order we intended?”




# Does Curricula give benefits?



- Standard1: mean performance over all 540 runs
- Standard2: split the 540 runs into 180 groups of 3; the maximum mean
- Standard3: the mean value of the top three values of all 540 runs

# Does Curricula give benefits?



- Marginal value of ordered learning
  - An artifact of the large search space
- No dependence on the three different orderings
  - In CIFAR10, the best **mean accuracy** is achieved via **random ordering**
  - In CIFAR100, the best single run has a random ordering

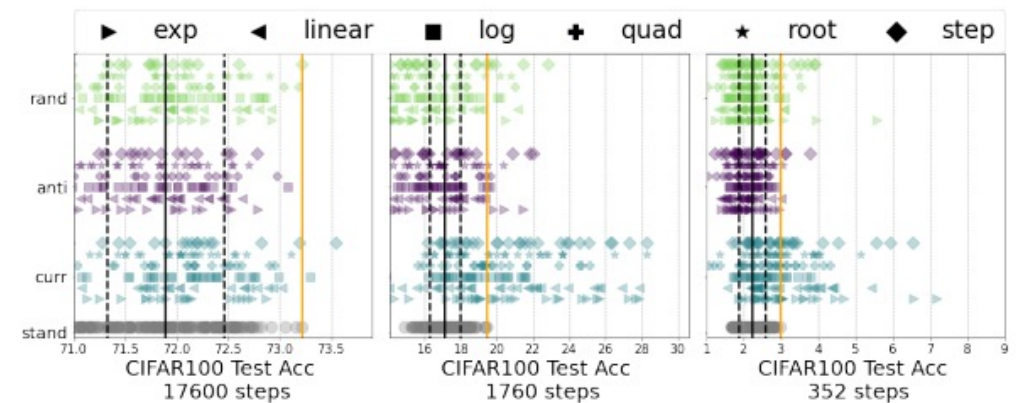

# Curricula For Short-Time Training And Noisy Data

- Many large-scale text models are trained using curricula
  - Data-rich setting(multiple epochs of training is not feasible)
  - Data is far less clean than standard image benchmarks.

| Dataset                 | Quantity<br>(tokens) | Weight in<br>training mix | Epochs elapsed when<br>training for 300B tokens |
|-------------------------|----------------------|---------------------------|---|
| Common Crawl (filtered) | 410 billion          | 60%                       | 0.44  |
| WebText2                | 19 billion           | 22%                       | 2.9   |
| Books1                  | 12 billion           | 8%                        | 1.9   |
| Books2                  | 55 billion           | 8%                        | 0.43  |
| Wikipedia               | 3 billion            | 3%                        | 3.4   |

**Table 2.2: Datasets used to train GPT-3.** “Weight in training mix” refers to the fraction of examples during training that are drawn from a given dataset, which we intentionally do not make proportional to the size of the dataset. As a result, when we train for 300 billion tokens, some datasets are seen up to 3.4 times during training while other datasets are seen less than once.

# Limited training time budget




- Curriculum learning can indeed improve performance when training time budget is decreased



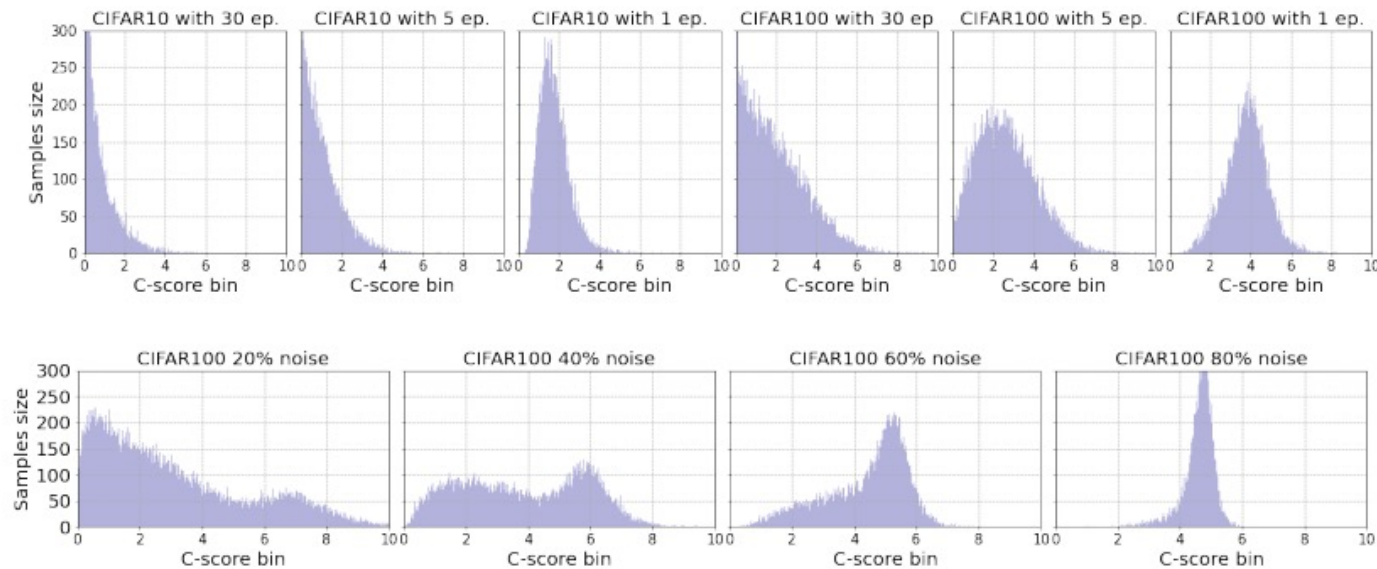
# Data With Noisy Labels

- Generating artificial label noise by randomly permuting the labels
  - Recompute the c-score



**Figure 7: Curriculum-learning helps when training with noisy labels.** Performance on CIFAR100 with the addition of 20%, 40%, 60% and 80% label noise shows robust benefits when using curricula.

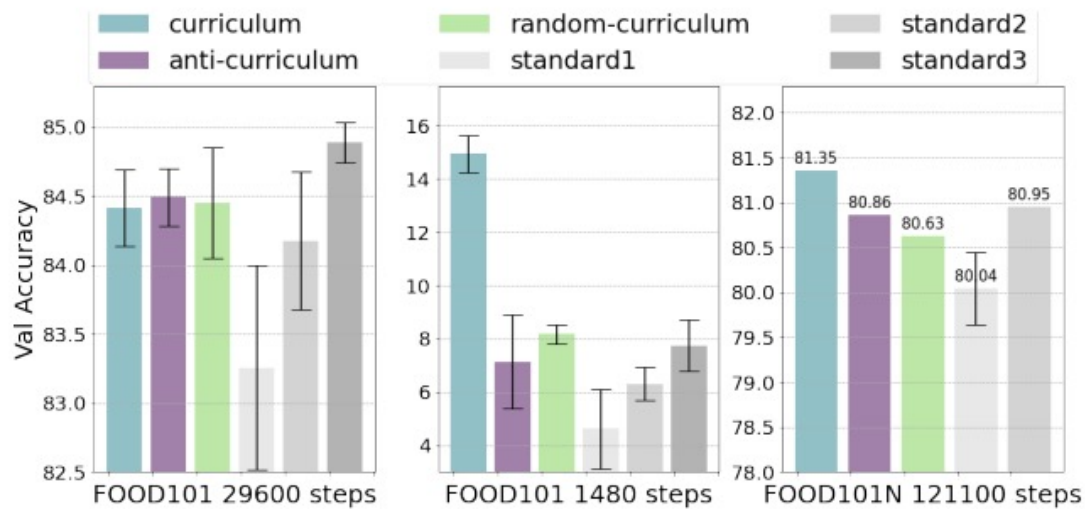
# Data With Noisy Labels



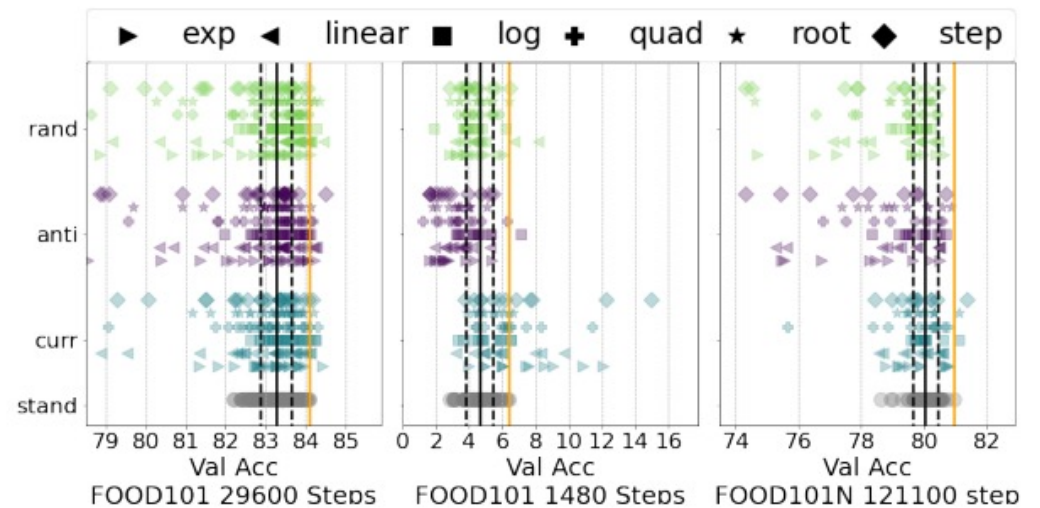
- Why does the label noise benefit from curricula?
  - A significant number of images concentrate around zero (clean CIFAR)



# Curricula In The Large Data Regime



(a)



(b)

- FOOD101(75,000 training/25,000 validation)
- FOOD101N(310, 000 training/25,000 validation) - Noisy

# Conclusion

- Implicit Curricula: Examples are learned **in a consistent order** (similar architecture, method)
  - And we can change this order by changing the order in which examples are presented during training
- Curricula achieve (almost) **no improvement** in the **standard setting**
- Curriculum learning **improves** over standard training when training **time is limited**
- Curricula improves over standard training in **noisy regime**